The idea. The progress.

Know your rights. Know your next step.

What are we building?

A legal-awareness assistant for India that helps people understand a situation, organize important facts, and prepare a practical next step. It does not replace a lawyer or promise an outcome.

01

Understand a problem

Possible issue, missing facts, sources and next steps.

02

Explain a document

Plain-language explanations of dates, amounts and clauses.

03

Prepare a draft

An editable request or complaint for careful human review.

WHERE WE ARE TODAY

Interactive concept demo is ready. Live AI, RAG and backend are still planned.

Project log

01 **Discovery completed**
Checked hardware, Python, Git and the existing local AI runtime.

02 **Product plan created**
Defined the audience, workflows, architecture and staged roadmap.

03 **Website demo built**
Problem, document, draft and trust/source views are available.

04 **Sample journeys added**
Prepared salary example and fictional letter; editable draft with copy/download.

05 **Core interactions checked**
Sample results, document answers, draft creation, copy and input rejection.

06 **Presentation material prepared**
Project Markdown, short pitch and mentor/judge questions are ready.

Delivery: local/offline HTML works. Online publication is not complete. Official resource links are included, but claim-level legal citations are not implemented.

How the finished app will work

Target architecture - only the browser concept demo exists today.

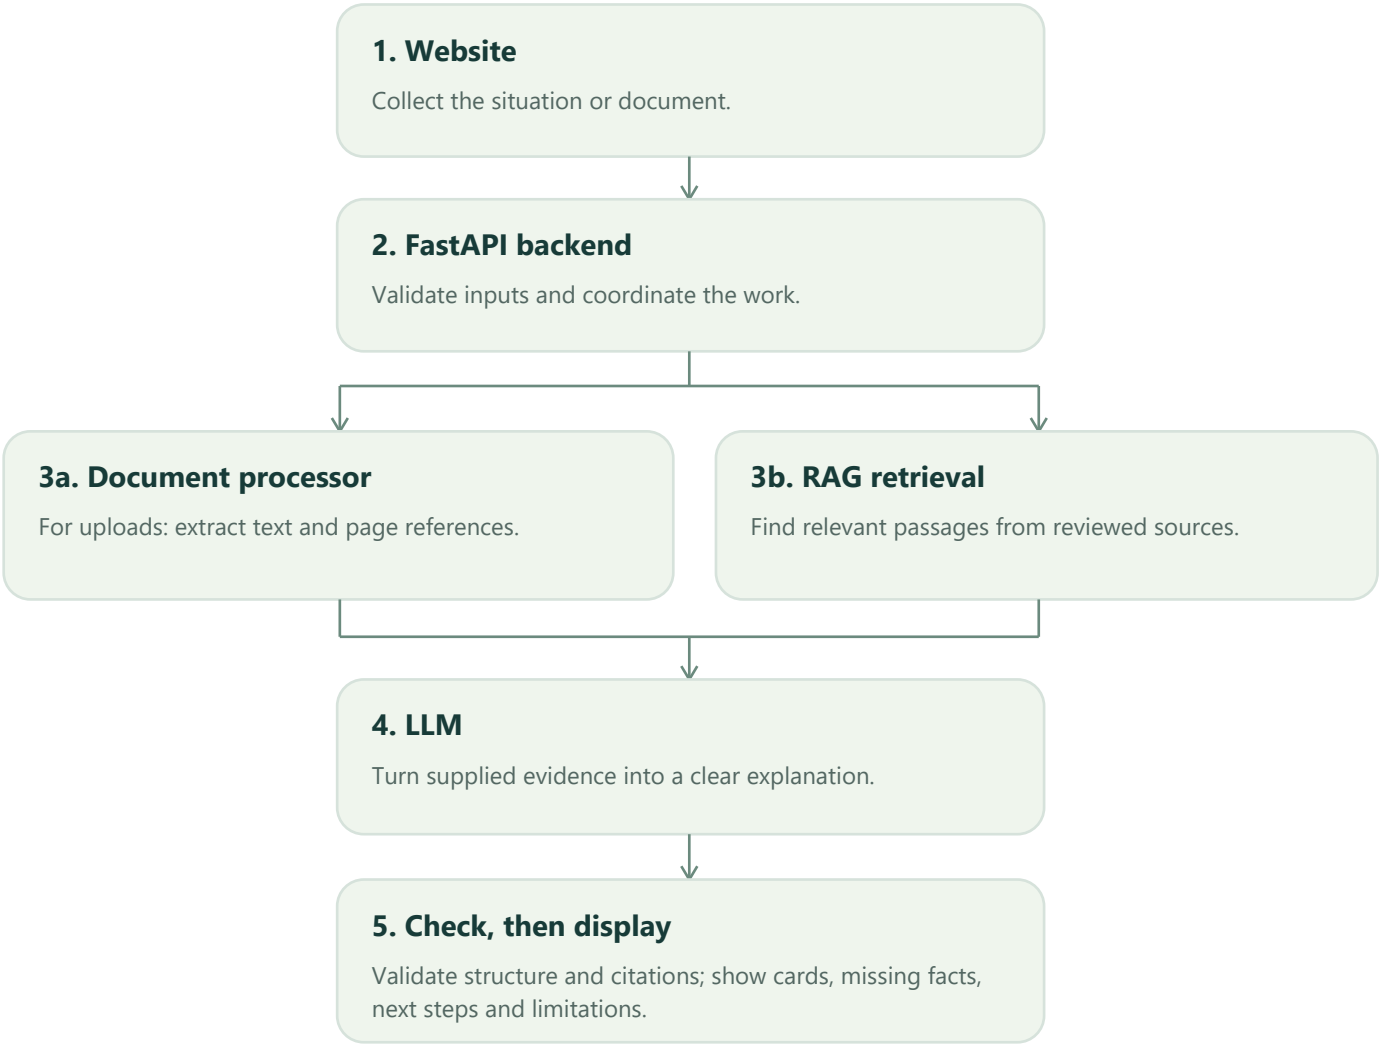

The components in everyday language

Frontend	HTML, CSS and JavaScript: the screens and buttons people use.
FastAPI	The Python server that connects the interface, AI and document tools.
LLM	Large language model: understands and writes text. Model choice is pending.
RAG	Retrieval-Augmented Generation: find evidence first, then supply it to the LLM.
SQLite	A small database for source metadata and review records; avoid storing personal inputs.

RAG preparation: official documents > small passages > embeddings (meaning as numbers) > similarity search. Retrieved passages keep their source references.

From demo to a working MVP

Build one reliable, evidence-backed journey before expanding features.

1

Connect the backend + LLM

Build FastAPI; test an installed local model before choosing a provider.

Pass when: valid requests work and timeouts/errors are handled.

2

Add sources + RAG + citations

Review a small official-source collection with dates and jurisdiction.

Pass when: important claims link to appropriate supporting passages.

3

Support real problem descriptions

Ask for missing facts and return structured, cautious answers.

Pass when: supported scenarios work and unsupported questions are declined.

4

Process documents + generate drafts

Add validated PDF/text uploads, document Q&A; and AI-assisted drafts.

Pass when: answers reference the document and unknown draft fields stay visible.

5

Test quality, safety and usability

Check grounding, injection, uploads, privacy, mobile layout and speed.

Pass when: core flows and failure cases pass; findings are recorded.

6

Package, deploy and rehearse

Set up the team repository, repeatable setup, deployment and demo backup.

Pass when: a teammate can run a complete 3-5 minute demonstration.

The strongest hackathon story

Real situation → **Visible evidence** → **Clear next step** → **Editable draft**

Show the source passage behind a claim, an honest response when evidence is missing, and a tested end-to-end flow. Add user feedback, qualified legal review where possible, and measured results. These strengthen the entry; no outcome is guaranteed.

Later: voice, OCR, regional languages, accounts and broad legal coverage. Current constraints: no hosted API access; local models are installed but not benchmarked. Privacy and citations reduce risk but cannot guarantee legal correctness.